


BOQ DATE DELIVERY PREDICTION (BOQ_DDP)

A RECURRENT NEURAL NETWORK FOR DELIVERY DATE PREDICTION



AGENDA

- Introduction
- Input data and preprocessing
- Architecture of RNN (LSTM)
- Training Process
- Results and their analysis
- Testing phase
- BoQ_DDP Web Application
- Conclusions



INTRODUCTION

1. Objective

- a. BoQ Delivery Date Prediction (BoQ_DDP) neural network, developed as part of a Proof-of-Concept (POC) to predict delivery dates for equipment items based on historical shipment data. Useful for Presales/PM/PEO

2. Dataset

- a. Source: Shipment delivery report 2024 and Q1 2025
- b. Size: 76,000 records
- c. Features: 14 variables
- Useful for training
 - Customer Request Date
 - Order In Take Date
 - Actual Ship Date
 - LeadTimeDays
 - IntakeLag

3. Model Goals


- a. Build an accurate and generalizable model for managing time sequences of events
- b. Analyze feature importance
- c. Visualize training behavior and performance metrics

4. Development

- a. Framework: Colab of Google, Tensorflow
- b. Language: Python 3.13 and 3.10
- c. Libraries: NumPy, Pandas
- d. Collaborator: ChatGPT 4o Plus

5. Bibliography and Inspiration:

- a) *Python and Machine Learning* A. Bellini, A. Guidi, McGraw-Hill. Study case: Predict the price of a stock, in our case Microsoft, starting from datasets available on Kaggle (<https://www.kaggle.com/dgawlik/nyse>)



INPUT DATA AND PREPROCESSING

INPUT DATA AND PREPROCESSING

- The dataset used to train and validate the BoQ_DDP neural network is the report of the shipment of IP product (EUR zone) covering 2024 and Q1 2025 (csv format), containing:
 - Over 76,000 samples (instances)
 - 14 features for each item
 - Integer label (LeadTimeDays) as target
- Features include information on:
 - Item_ID
 - Customer Request Date
 - Order In Take Date
 - Actual Ship Date
 - LeadTimeDays (as Actual Ship Date - Customer Request Date)
 - IntakeLag (as Order In Take Date - Customer Request Date)

INPUT DATA AND PREPROCESSING

- These phases of preprocessing are done **OUTSIDE** of neural network training script.
 - **LeadTimeDays** (as Actual Ship Date - Customer Request Date)
 - If LeadTimeDays is strictly less than -15, then the predicted LeadTimeDays must be equal to 0 (in practice, the shipment arrives on time)
 - If LeadTimeDays takes a value between -15 (inclusive) and zero (inclusive), then the predicted LeadTimeDays must be equal to 5 (in practice, the shipment arrives after a week equal to 5 days)
 - If LeadTimeDays takes a value strictly greater than zero and less than or equal to 225, then the predicted LeadTimeDays is equal to the actual LeadTimeDays.
 - If LeadTimeDays takes a value strictly greater than 225, then the predicted LeadTimeDays is equal to 225.
 - **IntakeLag** (as Order In Take Date - Customer Request Date)
 - If IntakeLag is negative, its value will be 1 (in line with the customer's request).
 - If IntakeLag is positive, its value will be 2 (not in line with the customer's request).
 - If IntakeLag is zero, its value will be 0 (somewhat in line with the customer's request).

INPUT DATA AND PREPROCESSING

- These phases of preprocessing are done **INSIDE** of neural network training script
 - **Convert the Item_ID column** (which contains product identification strings) into a numeric representation that the neural network can use as input. This is crucial because neural models cannot directly handle text or categorical data in string format. The saved *item_id_encoder.pkl* file is essential during the deployment phase to ensure consistency. The same Item_ID must always be encoded in the same way, even when used in production.
 - Let's assume that Item_ID contains: ['3HEI2345AA', '3HE67890BA', '3HEI2345AA']. After the encoding:
 - ['3HEI2345AA'] → 0
 - ['3HE67890BA'] → 1
 - The new column of Item_ID_enc is: [0, 1, 0]

INPUT DATA AND PREPROCESSING

- These phases of preprocessing are done **INSIDE** of neural network training script
 - **Sorting the dataset by:**
 - First for Item_ID_enc (i.e. groups all the same items together), then sort by Cust_Req_Ship_Date in ascending order within each group
 - Let's say you have the following unsorted dataset:

Item_ID_enc	Cust_Req_Ship_Date	LeadTimeDays
2	2025-07-18	12
1	2025-07-20	7
2	2025-06-15	10
1	2025-06-10	5
 - After the sorting:

Item_ID_enc	Cust_Req_Ship_Date	LeadTimeDays
1	2025-06-10	5
1	2025-07-20	7
2	2025-06-15	10
2	2025-07-18	12
 - This is essential for building the flowing timelines in the next steps, e.g. $[t-2, t-1] \rightarrow t$.

INPUT DATA AND PREPROCESSING

- **Normalization.** when training a neural network (or a machine learning algorithm in general), the input data must be scaled and normalized. Main reason: different feature scales can slow down, prevent or damage learning. **MinMax** normalization scales the feature LeadTimeDays in the range [0, 1]. Summarizing: stable learning, balanced gradients, better generalization, all features treated equally. The saved *leadtime_scaler.pkl* file is used during prediction to correctly scale on input and do the inverse on output.

$$x_{\text{scaled}} = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

$$x_{\text{original}} = x_{\text{scaled}} \times (x_{\max} - x_{\min}) + x_{\min}$$

- Suppose LeadTimeDays is: [4, 5, 6, 10]. After MinMaxScaler (Min: 4, Max: 10) → [0.0, 0.16, 0.33, 1.0] → Inverse Transform: [4.0, 5.0, 6.0, 10.0].
- In practice, *leadtime_scaler.pkl* does not contain the data, but the mathematical rules (parameters) used to transform it.
- If you don't use the same scaler, the input data would be normalized differently, and the model would receive out-of-scale values → incorrect predictions.

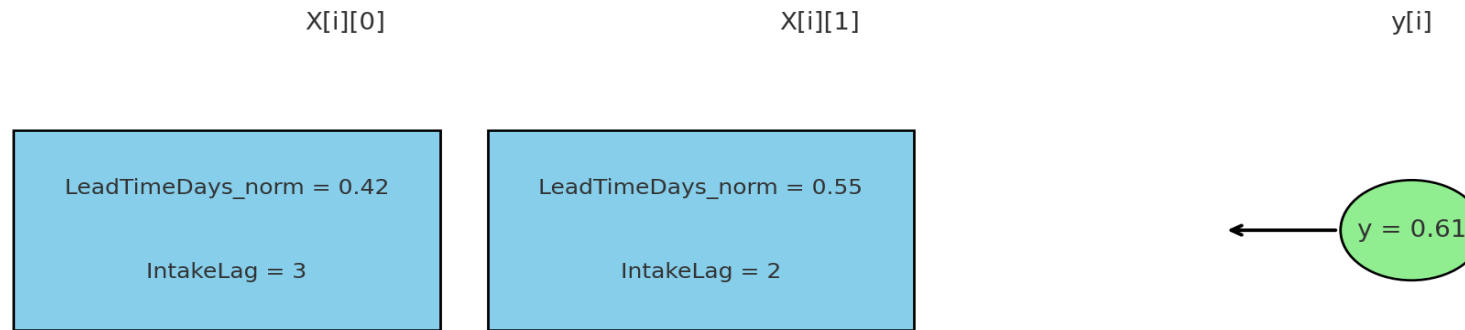
INPUT DATA AND PREPROCESSING

- Continuing
 - **Convert the DataFrame** into a time sequence dataset for each Item_ID, useful for training sequential models such as LSTM or GRU (sliding windows)
 - Suppose that for a certain Item_ID there are the following values:
 - LeadTimeDays_norm: [0.4, 0.5, 0.6, 0.8]
 - IntakeLag: [3, 2, 4, 1]
 - The complete sequence of features will be:
 - [[0.4, 3], [0.5, 2], [0.6, 4], [0.8, 1]]
 - Then with n_steps = 2 we create:
 - $X[0] = [[0.4, 3], [0.5, 2]] \rightarrow y[0] = 0.6$
 - $X[1] = [[0.5, 2], [0.6, 4]] \rightarrow y[1] = 0.8$

X is a 3D array: (number of sequences, n_steps, number of features)
y is a 1D array: target
corresponding to each sequence
n_steps = 2 means the model will use 2 previous rows to make a prediction

INPUT DATA AND PREPROCESSING

- This representation clearly shows how the data is packaged for training the LSTM/GRU model (see later).




- Here's a representation with a real-world numerical example:
 - The sequence $X[i]$ is composed of two time steps:
 - $X[i][0]$: LeadTimeDays_norm = 0.42, IntakeLag = 3
 - $X[i][1]$: LeadTimeDays_norm = 0.55, IntakeLag = 2
 - The target value $y[i]$ is 0.61, i.e., the value of LeadTimeDays_norm at the next time step.



INPUT DATA AND PREPROCESSING

- Continuing
 - **Splits the dataset** into a training set (70%) and a validation set (30%), based on a specified percentage. Training set is used for training phase; validation set is used for validation phase. Validation is used to verify that the training carried out with the training set does not generate overfitting.

LSTM



ARCHITECTURE OF
RNN (GRU+LSTM)

ARCHITECTURE OF RNN (GRU+LSTM)

- Why GRU+LSTM for Lead Time Prediction?
 - a. The problem is sequential/temporal in nature: lead time prediction relies on past chronological data per item (e.g., IntakeLag, historical LeadTimeDays).
 - b. Recurrent Neural Networks (RNNs) are suitable for sequences but suffer from long-term memory issues.
 - c. GRU and LSTM architectures enhance memory handling and temporal dependency learning.
 - d. The GRU + LSTM combination leverages:
 1. The **speed and simplicity** of GRU for short sequences.
 2. The **expressive power** of LSTM for learning deeper temporal patterns.
 - e. In this project, GRU+LSTM was chosen to:
 1. Handle lead time variability.
 2. Balance model complexity and predictive power.

ARCHITECTURE OF RNN (GRU+LSTM)

- How Does GRU+LSTM Work (Conceptually)?
 - a. **GRU (Gated Recurrent Unit):**
 1. Uses two gates: update and reset.
 2. Fewer parameters than LSTM; faster to train.
 - b. **LSTM (Long Short-Term Memory):**
 1. Uses three gates: input, forget, and output.
 2. Better suited for long-range temporal dependencies.
 - c. **In our model:**
 1. GRU processes the raw input and returns a sequence of representations.
 2. LSTM receives the sequence and learns a compact representation.
 3. Dense layers produce the final prediction.

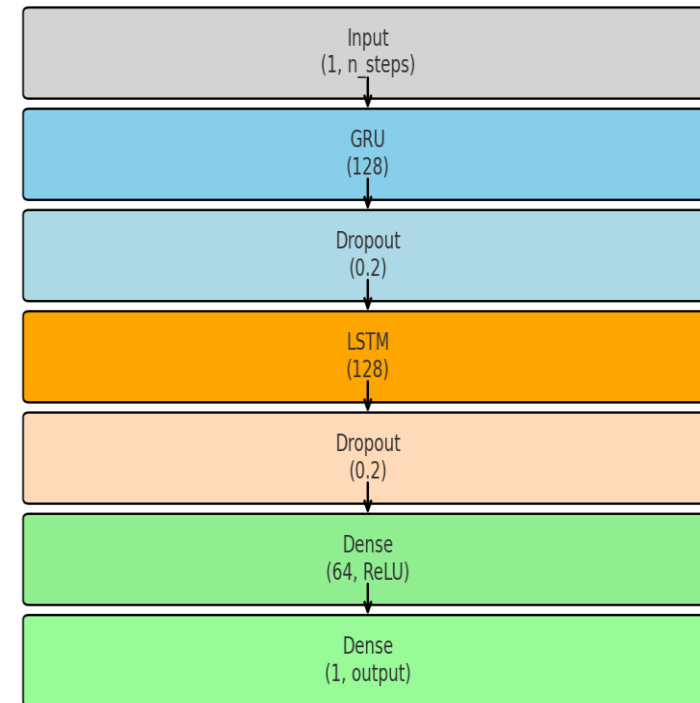
ARCHITECTURE OF RNN (GRU+LSTM)

- Key Characteristics of Our GRU+LSTM Model
 - **Sequential Architecture:**
 - GRU (128 units)
 - Dropout (0.2)
 - LSTM (128 units)
 - Dropout (0.2)
 - Dense (64, ReLU activation)
 - Dense (1, output)
 - **Design Rationale:**
 - GRU reduces early-stage overfitting.
 - LSTM captures deeper temporal sequences.
 - Dropout regularizes training.
 - Dense layer adds non-linearity and supports regression.

ARCHITECTURE OF RNN (GRU+LSTM)

- Simplified Graphical Representation of the Model
 - Block Diagram:
 1. Input: time sequences (LeadTimeDays_norm, IntakeLag)
 2. GRU Layer: 128 units → intermediate representation
 3. Dropout
 4. LSTM Layer: 128 units → deep temporal learning
 5. Dropout
 6. Dense (64, ReLU)
 7. Final Output: Dense (1)


Architettura del modello GRU + LSTM per Lead Time Prediction



ARCHITECTURE OF RNN (GRU+LSTM)

- Conclusion
 - The GRU+LSTM combination is a balanced solution for time-series prediction problems.
 - It allows:
 - Noise reduction in early-stage data (GRU)
 - Retention of useful sequence memory (LSTM)
 - Inference from historical component-based patterns.
 - Results indicate a good trade-off between accuracy and computational complexity.

ARTIFICIAL INTELLIGENCE



**MACHINE
LEARNING**



TRAINING
PROCESS,
RESULTS AND
ANALYSIS

TRAINING PROCESS, RESULTS AND ANALYSIS

- Model Compilation
 - Model compiled with the following configuration:
 - Loss function: Mean Squared Error (MSE), suitable for a regression problem
$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$
 - Optimizer: Adam (learning rate = 0.001), with an initial learning rate of 0.001. A balanced value between learning speed and stability
 - Metric: Mean Squared Error, again MSE is used for monitoring during training/validation
- EarlyStopping Callback
 - To prevent overfitting and unnecessary training:
 - Monitor: *val_mean_squared_error*, check the mean squared error on the validation data
 - Patience: 5 epochs, if there is no improvement for 5 consecutive epochs, training stops
 - Automatically restores best weights, upon completion, it automatically restores the epoch weights with the best *val_mse value*

TRAINING PROCESS, RESULTS AND ANALYSIS

- Training Phase
 - Model is trained using:
 - Training set: X_{train} , y_{train} , training data
 - Validation set: X_{test} , y_{test} , Validation performed at each epoch to monitor *val_mse*
 - Epochs: 50, Maximum number of epochs (complete cycles on the dataset)
 - Batch size: 32, the model updates the weights every 32 examples
 - Callback: EarlyStopping, activates the early stopping logic
- Training History
 - Output of *model.fit()* is stored in history:
 - Tracks epoch-wise MSE and val_MSE
 - Useful for plotting learning curves
 - Helps evaluate training quality

TRAINING PROCESS, RESULTS AND ANALYSIS

- Visual Flowchart



TRAINING PROCESS, RESULTS AND ANALYSIS


- Final Model Performance Metrics
 - MSE: 175.65 (days²) → Mean Squared Error
 - RMSE: 13.25 days → Root Mean Squared Error
 - MAE: **8.12 days** → Mean Absolute Error

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

- MAE (Mean Absolute Error)
 - *Answers to question: “On average, by how many days does the model miss the prediction?”*
- RMSE (Root Mean Squared Error)
 - *Answers to question : “How much does the model miss on average, giving more weight to large errors?”*
- MSE (Mean Squared Error)
 - *Answers to question : “What is the average variance of the errors (in days²) made by the model?” not immediate, loss function*

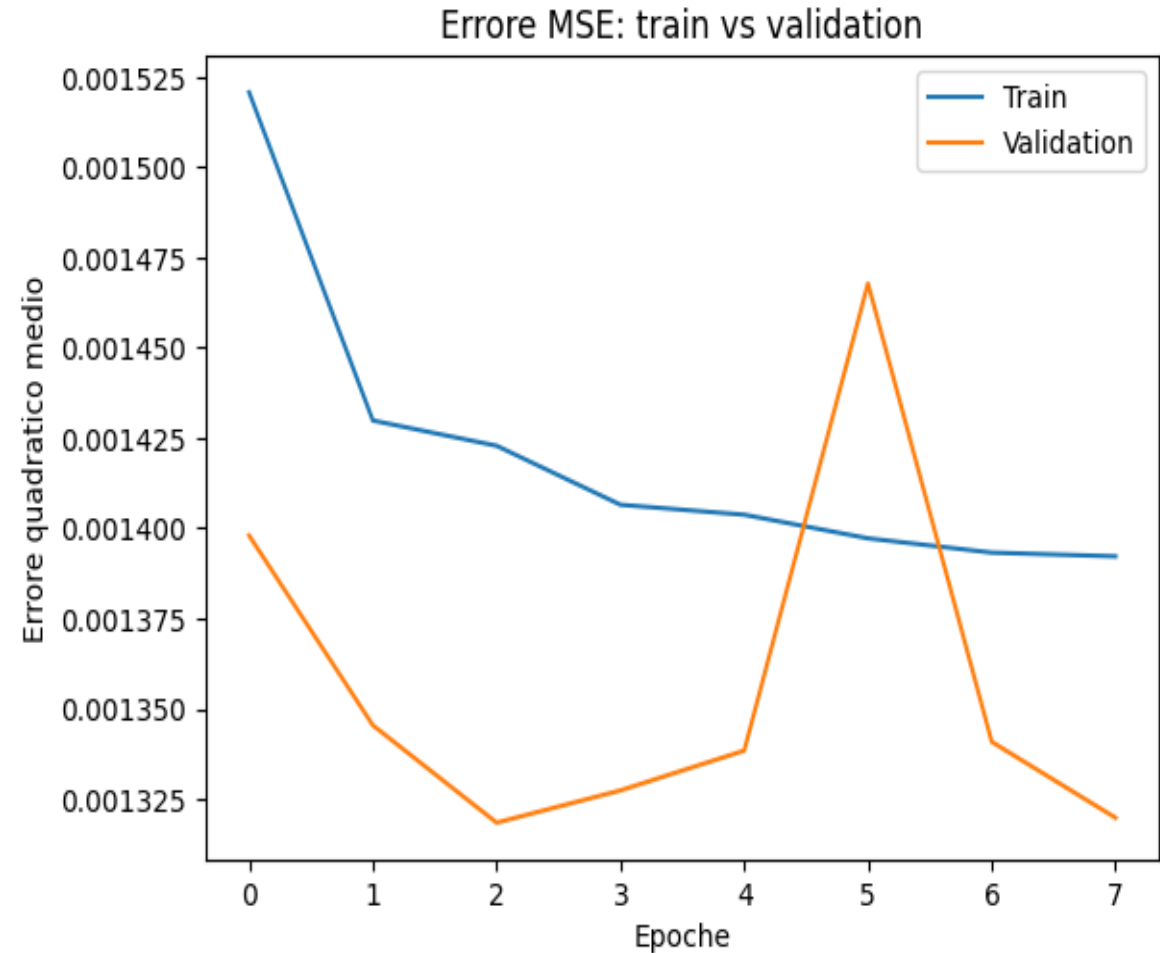
TRAINING PROCESS, RESULTS AND ANALYSIS

- LeadTimeDays: Actual vs Predicted
 - The model **closely follows the overall trend** of actual lead times
 - Predicted values are **clipped** to prevent unrealistic outliers
 - Slight underestimation in high-peak cases, but performance is solid overall
- Analysis: Actual vs Predicted
 - Strong accuracy for **low and medium** lead times
 - Slight underestimation in **high peaks**, handled with clipping
 - Best predictions occur in **densely represented regions** of the dataset



TRAINING PROCESS, RESULTS AND ANALYSIS

- Mean Squared Error: Train vs Validation
 - MSE steadily decreases for both training and validation sets
 - No signs of overfitting
 - EarlyStopping triggered after 7 epochs with best weights restored





TRAINING PROCESS, RESULTS AND ANALYSIS

- Training Conclusions
 - Training process is stable and well-regularized
 - EarlyStopping stopped at the optimal moment (epoch 7)
 - The model has robust predictive capacity, with $MAE < 10$ days
 - Ready for production usage

AI TESTING



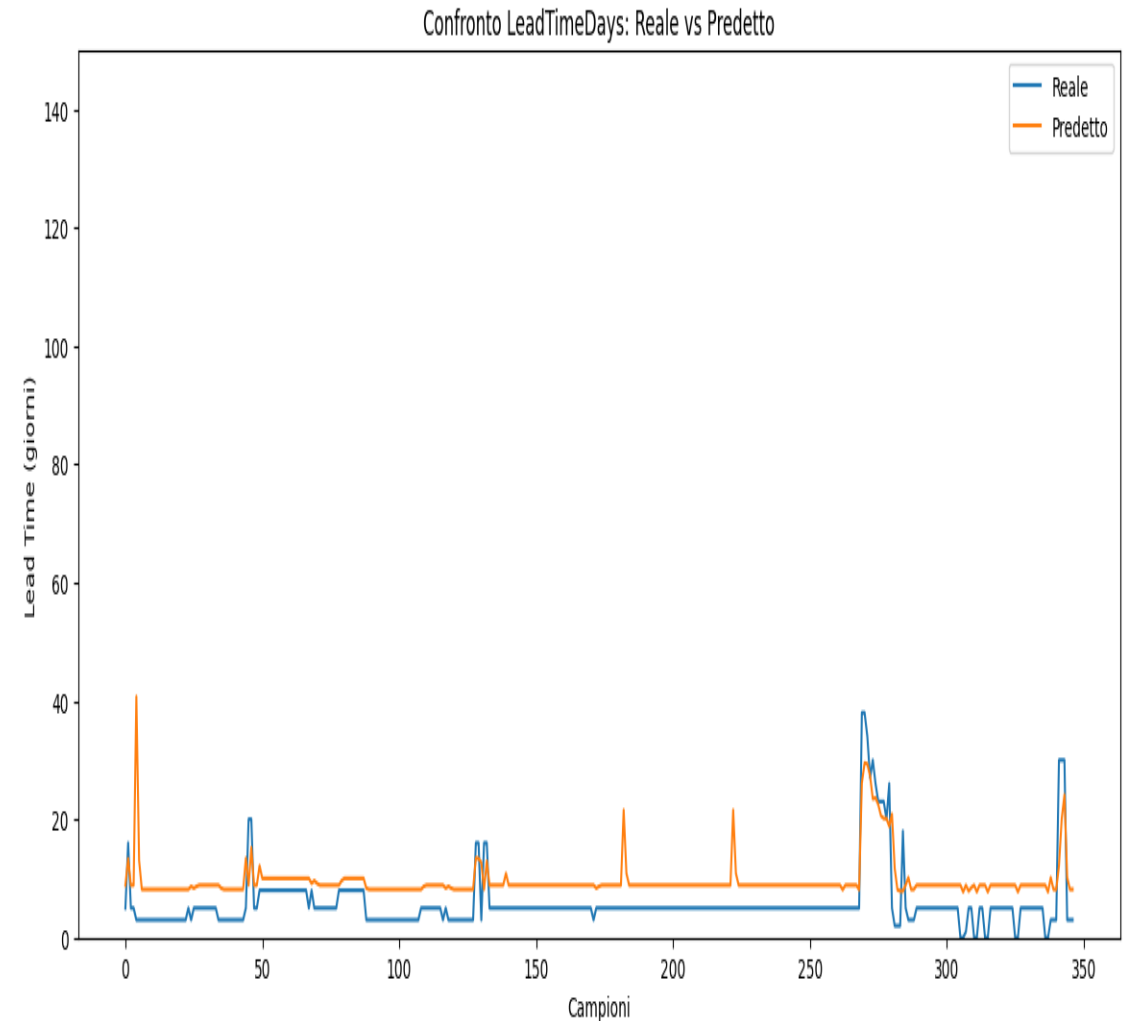
TESTING PHASE

TESTING PHASE

- Script Workflow Overview
- The test script performs the following sequential steps:
 - Load LSTM-GRU trained model (.keras)
 - Upload test dataset with delivery history (about 470 items)
 - Encode categorical variable Item_ID
 - Normalize LeadTimeDays using saved scaler
 - Filter items with at least 3 records
 - Generate time sequences with IntakeLag + LeadTimeDays_norm
 - Predict normalized LeadTimeDays using the model
 - Reverse normalization to obtain results in days
 - Plot predictions vs actuals and calculate error metrics


TESTING PHASE

- Predicted vs Actual Lead Time
- Graphical comparison of predicted vs actual delivery lead time (in days):
 - Blue line: Ground truth from dataset
 - Orange line: Model predictions
 - Observations:
 - Model is smooth, avoids noise
 - Accurate in normal ranges (5–15 days)
 - Underestimates delivery delays (spikes >30d)



TESTING PHASE

- Test Metrics Summary
- Final performance metrics on the test dataset:
 - MAE (Mean Absolute Error): 6.51 days → Average error, robust to outliers. Good average error tolerance (**MAE < 7d**)
 - RMSE (Root Mean Squared Error): 7.72 days → Penalizes large deviations
 - MSE (Mean Squared Error): 59.66 days² → Measures overall error dispersion





BOQ DATE
DELIVERY
PREDICTION WEB
APPLICATION

BOQ DATE DELIVERY PREDICTION WEB APPLICATION

- A web application developed with Streamlit for predicting the delivery date of a BoQ using a pre-trained LSTM-GRU neural network.
- Key functionalities:
 - Upload of an Excel file with Item_ID entries to be evaluated
 - Input of two dates: Customer Request Date (CRD) and Order Intake Date (OID)
 - Calculation of IntakeLag and sequence generation for each item
 - Prediction of Lead Time (in days) using the neural model
 - Automatic calculation of the predicted delivery date (excluding weekends)
 - Tabular display of predicted dates with confidence level
 - Informative messages about model performance and options to reset or exit the app

Demo





CONCLUSION

CONCLUSION

- Benefit
 - Captures time-based patterns with LSTM-GRU
 - Good accuracy: MAE \approx 8.12 days, RMSE \approx 13.25 days
 - Works well with just 2 input features
 - Generalizes across many Item_IDs
- Future Improvements
 - Add more features (e.g., category, forecast type)
 - Handle extreme delays more accurately
 - Cluster similar items for specialized models
 - Set up periodic retraining with new data



THANK YOU

- Corrado Vaccaro
- Senior Customer Program Manager, PMP®
- corrado.vaccaro@nokia.com